
Playing Through an AI Incident

Sofia G. Gallego
Cosmoventures, Paris, France

Abstract

Little Agent Lab is a small browser game about how AI agents get permissions, and about trusting a safeguard for what it actually does rather than what it's called. You put together a small team of agents, decide who can access what, and watch what really happens when the plan runs. Three short missions build on each other, ending with one where you have to run the plan and read the result to find out which of two guards is the one actually watching the gate. The idea is early awareness: giving people, especially kids, a hands-on feel for ideas like oversight and verification, playfully and early. This is a small, early version of the idea that I'd like to keep developing.

1 What it is

Little Agent Lab is a browser game and a small, deterministic simulation you run directly, with results you can watch happen. You pick a builder agent, add up to two helper agents, set a gate that controls outside access, and run the plan. What happens plays out as a short, literal trace of events, step by step.

Three missions, unlocked in sequence:

- **Weather Lantern.** Two strategies work here: give the builder a helper whose route is separate from the private item, or open the gate and pair it with an inspector that checks what leaves. A first taste of capability separation.
- **Garden Delivery.** The winning move is to take the builder's own route out through the gate and pair it with a real check.
- **Message Relay.** Two helpers offer to guard the gate. Running the plan and reading the trace tells you which one actually does the job. This is the core idea of the whole project: a safeguard is worth exactly what it does.

2 Why

Most explanations of AI oversight are prose aimed at people who already work in the field. This project is aimed the other way: at a general audience, and especially at kids, who will grow up as the users of these systems. The bet is that a five-minute playable moment — picking the wrong guard and watching what happens — builds an intuition that a paragraph doesn't. That intuition, at scale, is a small but real input into how the next generation of people reasons about AI systems around them.

3 What's next

All three missions' logic is covered by automated tests (27, all passing) and checked in a real browser. Next: more missions, a version where the fake safeguard's tell changes each time so the lesson general-

izes instead of being memorized, and a playtest with the age group it's meant for.

Limitations and Dual-Use Considerations

Limitations. No human-subject testing has been done; every usability decision so far came from one non-blinded evaluator. Three missions is enough to show the mechanic, not enough to know if the lesson generalizes. The game's five-line trace is far more legible than a real system's actual safeguards, and a player could walk away overconfident about how easy real verification is. The implementation was AI-assisted (Claude Code), which I'm disclosing here per the sprint's guidance on describing contributions accurately.

Dual-use. Risk is low. Every agent name, destination, and payload in the game is fictional; nothing represents a real API, protocol, or system. The one mechanic that resembles a "bypass" (Message Relay's fake guard) is a fixed, pre-authored branch chosen from two labeled cards, not a generative or discoverable exploit — there is nothing here a player could apply against a real system. No model was jailbroken and no organization's real systems were touched in building this.

Project source: <https://github.com/nefinia/little-agent-lab> (public, playable via the repo's GitHub Pages workflow).