
Playing Through an AI Incident

Sofia G. Gallego
Cosmoventures, Paris, France

Abstract

Little Agent Lab is a small browser game where you build a team of AI agents to finish a task while keeping something private safe. You choose who's on the team, decide how much access each one gets, and then run the plan and watch what actually happens. There are three missions so far, and the last one has two guard characters where only one of them really does its job – you find that out by running the plan and reading the trace, not by reading a description. I made this because most explanations of AI oversight are written for people who already understand the topic, and I wanted something a kid, or anyone new to it, could just play with and get a feel for. It's still small and early, and I want to keep building on it.

1 What it is

It's a browser game, nothing to install. You pick a builder agent, optionally add one or two helpers, decide whether the builder is allowed to leave the workshop through a gate, and press run. Then you watch a short trace of what each agent actually did, one line at a time.

There are three missions right now:

- **Weather Lantern.** You can win this one two ways: give the builder a helper who never touches the private item, or open the gate and add an inspector who checks what goes out.
- **Garden Delivery.** The builder has to go through the gate itself this time, so the only way through is pairing that with a real inspector.
- **Message Relay.** Two helpers offer to guard the gate, but only one of them actually stops anything. You won't know which until you run the plan and see what happened.

2 Why

Most writing about AI oversight assumes you already work in the field. I wanted to build something for people who don't, especially kids, since they're the ones who'll grow up using these systems. Playing through a wrong guess and watching what happens teaches something a paragraph of explanation usually doesn't.

3 What's next

All three missions have automated tests behind them (27, all passing), and I've checked the whole thing in a real browser. Next I want to add more missions, make the fake guard change from run to run so people can't just memorize which one is real, and actually get some kids to play it and see what happens.

Limitations and Dual-Use Considerations

Limitations. I haven't tested this with real players yet, so everything about how it plays comes from my own judgment. Three missions show the idea but aren't enough to know if the lesson holds up more broadly. The trace in the game is also a lot easier to read than any real safeguard would be, so someone could come away thinking real verification is quicker and simpler than it actually is. I used Claude Code throughout to build this, which I'm noting here since the sprint asks for that to be clear.

Dual-use. I think the risk here is low. Every name, destination, and payload in the game is made up, and none of it maps to a real API or system. The closest thing to a "bypass" is Message Relay's fake guard, but it's just one of two fixed cards to pick from, not something a player discovers or generates, so there's nothing to actually apply against a real system. No model was jailbroken and no real systems were involved in making this.

Project source: <https://github.com/nefinia/little-agent-lab> (public, playable via the repo's GitHub Pages workflow).